

# Full Implementation Architecture

## Industry-Grade AI Commerce Assistant (RAG + Tools + LCEL + OpenAI)

This document defines a **production-ready folder structure**, modular architecture, service boundaries, and component responsibilities for a scalable AI Commerce Assistant.

This system integrates:

- OpenAI LLMs
- OpenAI embeddings
- LangChain LCEL pipelines
- Chroma VectorDB
- RAG retrieval
- Tool calling (Weather + Serper + Utilities)
- Hybrid fallback strategy
- Clean separation of concerns

---

## 1. High-Level System Architecture


Client (Web / Mobile / Chat)



API Gateway (FastAPI)



Application Layer (Orchestrator)



---

## 2. Clean Architecture Layers

### 2.1 Presentation Layer

- FastAPI endpoints
- Request validation
- Response formatting

### 2.2 Application Layer

- Orchestration logic
- Intent routing
- Pipeline execution
- Tool loop handling

## 2.3 Domain Layer

- Recommendation rules
- Retrieval policies
- Ranking logic
- Business constraints

## 2.4 Infrastructure Layer

- VectorDB (Chroma)
  - OpenAI API integration
  - Serper API integration
  - Weather API integration
- 

## 3. Production Folder Structure

ai-commerce-assistant/

```
|  
|   └── app/  
|       |   └── main.py  
|       |   └── config.py
```

```
|   └── dependencies.py  
|  
|   |  
|   └── api/  
|       |   └── routes.py  
|       |   └── schemas.py  
|       └── middleware.py  
|  
|   |  
|   └── core/  
|       |   └── orchestrator.py  
|       |   └── router.py  
|       |   └── intent_classifier.py  
|       └── response_builder.py  
|  
|   |  
|   └── rag/  
|       |   └── ingestion.py  
|       |   └── chunking.py  
|       |   └── embeddings.py  
|       |   └── vector_store.py  
|       |   └── retriever.py  
|       └── rag_pipeline.py  
|  
|   |  
|   └── recommendations/  
|       |   └── recommender.py  
|       |   └── ranking.py  
|       └── personalization.py
```

```
| |
|   |   tools/
|   |   |   weather.py
|   |   |   serper_search.py
|   |   |   calculator.py
|   |   |   interest.py
|   |   |   tool_registry.py
|   |
|   |   llm/
|   |   |   openai_client.py
|   |   |   prompts.py
|   |   |   lcel_pipelines.py
|   |   |   tool_executor.py
|   |
|   |   database/
|   |   |   product_loader.py
|   |   |   metadata_filters.py
|   |
|   |   monitoring/
|   |   |   logging.py
|   |   |   metrics.py
|   |   |   tracing.py
|   |
|   |   utils/
|   |   |   helpers.py
```

```
|   └── constants.py  
|  
|  
└── data/  
    |   └── raw_catalog.csv  
    |   └── chroma_db/  
    |  
    |  
    └── tests/  
        |   ├── test_rag.py  
        |   ├── test_tools.py  
        |   ├── test_router.py  
        |   └── test_api.py  
        |  
        |  
        └── scripts/  
            |   ├── ingest_catalog.py  
            |   └── rebuild_embeddings.py  
            |  
            |  
            └── docker/  
                |   ├── Dockerfile  
                |   └── docker-compose.yml  
                |  
                |  
                └── requirements.txt  
                └── .env  
└── README.md
```

---

## 4. Component-Level Design

---

### 4.1 API Layer (`api/`)

#### `routes.py`

Responsibilities:

- Expose `/chat`
- Accept user query
- Call orchestrator
- Return structured response

POST /chat

```
{  
    "user_id": "123",  
    "query": "Recommend waterproof shoes under 4000"  
}
```

---

### 4.2 Orchestration Layer (`core/`)

#### `orchestrator.py`

Responsibilities:

- Manage state
- Invoke intent classifier

- Route to appropriate pipeline
- Execute tool loop if required

Pseudo-flow:

```
def handle_request(query):

    intent = classify_intent(query)

    if intent == "product_search":

        return rag_pipeline.run(query)

    if intent == "recommendation":

        return recommender.run(query)

    if intent == "weather":

        return tool_executor.run(query)

    return fallback_handler.run(query)
```

---

## 4.3 RAG Module (**rag/**)

### **ingestion.py**

- Load product catalog
- Normalize text

### **chunking.py**

- RecursiveCharacterTextSplitter
- Configurable chunk size

## **embeddings.py**

- OpenAI embedding wrapper

## **vector\_store.py**

- Chroma initialization
- Persistence handling

## **retriever.py**

- Similarity search
- Metadata filtering

## **rag\_pipeline.py**

LCEL-based retrieval pipeline.

```
rag_chain = (
    {
        "context": retriever | format_docs,
        "question": RunnablePassthrough()
    }
    | prompt
    | llm
    | StrOutputParser()
)
```

---

## 4.4 Recommendation Engine (recommendations/)

### recommender.py

- Applies business rules
- Calls retriever
- Applies ranking

### ranking.py

- Score normalization
- Hybrid ranking

### personalization.py

- User embedding generation
  - Similarity scoring
- 

## 4.5 Tool Layer (tools/)

Each tool is isolated.

Example:

### weather.py

- OpenWeather API call

### `serper_search.py`

- Serper API wrapper

### `tool_registry.py`

```
TOOL_MAP = {  
    "get_weather": get_weather,  
    "web_search": web_search,  
    "calculate_interest": calculate_interest  
}
```

This avoids hardcoding in orchestrator.

---

## 4.6 LLM Layer (**llm/**)

### `openai_client.py`

- Centralized OpenAI config
- Temperature control
- Model selection

### `prompts.py`

- All prompt templates
- Strict grounding rules

### `lcel_pipelines.py`

- RAG pipelines

- Tool pipelines
- Hybrid pipelines

## **tool\_executor.py**

- Automatic tool loop
  - Safety validation
- 

## **4.7 Monitoring Layer**

- Structured logging
- Tool invocation logs
- Retrieval score logging
- Token usage tracking
- Latency measurement

Enterprise systems require observability.

---

## **5. Deployment Architecture**

Load Balancer



FastAPI Service (Docker)



ChromaDB (Persistent Volume)

↓

External APIs:

- OpenAI
  - Serper
  - OpenWeather
- 

## 6. CI/CD Strategy

- Unit tests for tools
  - Integration tests for RAG
  - LLM response validation
  - Retrieval quality evaluation
  - Dockerized builds
  - Environment-based config
- 

## 7. Scaling Considerations

- Stateless API servers
- Shared vector database
- Caching frequent queries
- Tool result caching
- Rate limiting
- Retry logic for APIs

---

## 8. Security Considerations

- API key management via environment variables
  - Tool execution whitelisting
  - Input validation
  - Prompt injection guardrails
  - Metadata filtering enforcement
- 

## 9. Extension Roadmap

Future expansion:

- LangGraph stateful orchestration
  - Multi-agent architecture
  - Cross-encoder reranking
  - Evaluation framework
  - Feedback-driven learning loop
- 

## 10. Final Summary

This folder structure and architecture:

- Enforces separation of concerns

- Enables modular scalability
- Supports tool calling and RAG
- Uses LCEL pipelines
- Supports hybrid retrieval
- Is deployable in enterprise environments
- Allows observability and governance